

## ECE 361E: Machine Learning and Data Analytics for Edge AI

### HW1 Assigned: Jan 17 DUE: Jan 26 (11:59:59pm CST)

Work in groups of two students. At the end of the PDF file, insert a paragraph where you describe each member's contribution and two valuable things you learned from this homework.

Only one submission per group is required.

## Introduction

This assignment is meant to be an introduction to training and testing Machine Learning (ML) models on the [MNIST dataset](#)<sup>1</sup> in the Cloud. Specifically, you will be working with [PyTorch](#), one of the most popular frameworks for developing ML models. By working on this assignment, you will learn:

- The basics of ML techniques (e.g., training and testing a model, handling overfitting, comparing and contrasting optimizers, evaluating the complexity of a model and fine-tuning the hyperparameter of a model);
- How to visualize various metrics and parameters and their impact on different design decisions;
- How to discuss the significance of your experimental results.

## Problem 1 [20p]: Logistic Regression (using TACC Machines)

**Question 1: [5p]** Modify *starter.py* by specifying the target device for your model `model = model.to(torch.device('cuda'))` to train the model on GPU. Run *starter.py* on Maverick2 (see **Appendix A1.1** for environment setup, **Appendix A1.2** to run your code, and **Appendix A1.3** for more info).

**Question 2: [6p]** Draw a plot showing the variation of *training loss* and the *test loss* of your model for each *epoch* (this is the *loss plot*). Draw a second plot with the *training* and *test accuracies* of your model for each *epoch* (this is the *accuracy plot*).

**Question 3: [9p]** Complete *Table 1*:

Table 1

Training accuracy [%]	Testing accuracy [%]	Total time for training [s]	Total time for inference [s]	Average time for inference per image [ms]	GPU memory during training [MB]

## Problem 2 [40p]: Overfitting, Dropout and Normalization

**Question 1: [8p]** Run the code from *simpleFC.py* as is. Draw the loss and accuracy plots. Does this model overfit? Explain.

**Question 2: [14p]** In the `__init__()` method from the *SimpleFC* class, define once `nn.Dropout(probability)` and apply this new operation in the `forward()` method after every layer of the model except the last one. Run four experiments using the values [0.0, 0.2, 0.5, 0.8] as probabilities for dropout. Draw one loss plot for each experiment. What do you observe from these plots? Which dropout probability gives the best/worst results (the best results have no overfitting, i.e., almost equal train and test loss values)? Explain.

**Question 3: [18p].** Complete *Table 2* by running the same code as in **Problem 2, Question 2**, but using only the best dropout rate. In *Table 2*, replace X with the dropout rate which had the best results in

<sup>1</sup> [Modified National Institute of Standards and Technology database \(MNIST\)](#)

**Problem 2, Question 2.** For the row with “+ norm”, keep the same dropout probability and normalization to both training and testing datasets. Compare and contrast these normalized and unnormalized experiments and explain the differences.

Table 2

Dropout	Training accuracy [%]	Testing accuracy [%]	Total time for training [s]	First epoch when the model reaches 96% training accuracy
X				
X + norm				

### Problem 3 [40p + 10Bp]: CNNs and Model Complexity

**Question 1: [20p]** Implement the normalized MNIST dataset from **Problem 2, Question 3** in *simpleCNN.py*. Run the code and complete **Table 3**. Is there a difference between the estimated (total) size of the model from *torchsummary* and the saved version of the model? Explain.

Table 3

Model name	MACs	FLOPs	# parameters	Model size [MB]	Saved model size [KB]
SimpleCNN					

**Question 2: [20p]** Create your own CNN with at least two convolutional layers and train it on the normalized MNIST dataset. Draw the loss and accuracy plots. Further extend **Table 3** with a new row with your model name and the new results.

**BONUS Question 3: [10p]** Explain your choice for the model architecture. Does your model overfit? How does your proposed model compare against the other models from this homework? Does it have more/less parameters? Is it smaller/bigger in size? Would you use your model or any of the others? Why or why not?

### Submission Instructions

Include your solutions into a single zip file named <Group#>.zip. The zip file should contain:

1. A single PDF file containing all your results and discussions.
2. A *readme.txt* file explaining all your items in the zip file.
3. Your code files, named suggestively (e.g., **p1\_q1.py** for **Problem 1 Question 1** code).

***Good luck!***